

Data Card

Input vs. Output	- What are users inputting (for example, audio, text, images)? - What are the outputs or labels your model will predict or generate?	Inputs: Raw text sequences written in various target African languages. Outputs: Sub-word tokens (string fragments) and their corresponding numerical token IDs mapped directly from the trained tokenizer's vocabulary integer index.
Classification problem	- What is the classification task, if any? - What categories or classes are you working with, and why?	Tokenization itself is a sequence segmentation and vocabulary lookup task rather than a traditional multi-class classification problem. The model transforms input character strings into discrete sequence representations (ordered integer lists) to optimize computational efficiency and alignment for downstream language model processing.
Labelling	- Where do your labels come from—domain experts, community members, crowd workers, AI-generated sources? - How will you check for consistency and accuracy?	Our text corpora and ground-truth tokens are aggregated from domain experts, native-speaking community contributors, and established public GitHub open-source repositories dedicated to African NLP. To verify accuracy and consistence, especially for languages outside our team's direct fluency. We perform systematic audits of repository documentation, community discussion threads, issue trackers, and peer-reviewed benchmark sources.
Data creation	- How will you collect your data? - Will it come from direct sources (like interviews or recordings), or will you synthesize data in low-resource settings (e.g., for under-resourced languages)?	We collect primary training corpora by curating existing open-source datasets from verified GitHub repositories. These sources contain pre-collected text spans spanning diverse media, including news archives, religious translations (e.g., Bible corpora), and literature. Because sufficient natural language text exists across our target language families, synthetic data generation is currently unnecessary
Web collection ethics	- Do you plan to source data from the web? - If yes, which sites or platforms? - What tools will you use - AI tools, human workers, or both?	We primarily rely on vetted, open-license GitHub repositories rather than automated web scraping. In instances where web data collection becomes necessary, we enforce strict compliance checks regarding licensing, terms of service, consent, and privacy rights before ingestion. We audit copyright metadata for every data source to ensure ethical alignment and non-infringement.

	Have you addressed issues of consent, copyright, and privacy?	
Bias and representation	- Who might be left out of your dataset? - How will you ensure that the dataset reflects the diversity of the community and includes underrepresented voices?	Limitations & Bias: Because our training corpora primarily derive from religious texts and news archives, out-of-domain vocabulary (e.g., informal conversational speech, medical, or legal terms) may be sub-optimally tokenized. Representation Strategy: To maximize representation under computational constraints, we sample representative languages across major African language families (and sub-families). Furthermore, we employ temperature-based data sampling (adjusting the sampling α-parameter during tokenizer training) to prevent high-resource languages from dominating the vocabulary allocation.
Values and community alignment	- How does your dataset design reflect the values you identified in 01 Build Your Own Small Language Model - How will your community benefit from this dataset, and how will you ensure accountability to them?	Value Alignment & Preservation: Our methodology embeds inclusivity, cultural preservation, and reliability. We preserve linguistic integrity by preserving native orthography, diacritical marks, and tonal accents. Technical reliability is verified through round-trip encoding/decoding fidelity checks. Community Impact & Accountability: Open-sourcing this efficient tokenizer lowers compute barriers for researchers building downstream African language models. We ensure accountability by maintaining open issue channels, accepting community feedback, releasing full documentation, and iteratively patching model defects.